
Supplementary Material

Anonymous Author(s)

Affiliation

Address

email

1 Proofs from section: Model and problem formulation

1.1

Lemma 1 Assuming a second order Taylor approximation, the quantization sensitivity $\Gamma(X, \epsilon)$ satisfies the following equation:

$$\Gamma(X, \epsilon) = \left| \frac{\partial^2 \text{MSE}(X, \Delta = \tilde{\Delta})}{\partial^2 \Delta} \cdot \frac{\epsilon^2}{2} \right|. \quad (1)$$

Proof: Let Δ' be a quantization step with similar size to $\tilde{\Delta}$ so that $|\Delta' - \tilde{\Delta}| = \epsilon$. Using a second order Taylor expansion, we approximate $\text{MSE}(X, \Delta')$ around $\tilde{\Delta}$ as follows:

$$\begin{aligned} \text{MSE}(X, \Delta') &= \text{MSE}(X, \tilde{\Delta}) + \frac{\partial \text{MSE}(X, \Delta = \tilde{\Delta})}{\partial \Delta} (\Delta' - \tilde{\Delta}) \\ &\quad + \frac{1}{2} \cdot \frac{\partial^2 \text{MSE}(X, \Delta = \tilde{\Delta})}{\partial^2 \Delta} (\Delta' - \tilde{\Delta})^2 + O(\Delta' - \tilde{\Delta})^3. \end{aligned} \quad (2)$$

Since $\tilde{\Delta}$ is the optimal quantization step for $\text{MSE}(X, \Delta)$, we have that $\frac{\partial \text{mse}(X, \Delta = \tilde{\Delta})}{\partial \Delta} = 0$. In addition, by ignoring order terms higher than two, we can re-write Equation (2) as follows:

$$\text{MSE}(X, \Delta') - \text{MSE}(X, \tilde{\Delta}) = \frac{1}{2} \cdot \frac{\partial^2 \text{MSE}(X, \Delta = \tilde{\Delta})}{\partial^2 \Delta} (\Delta' - \tilde{\Delta})^2 = \frac{\partial^2 \text{MSE}(X, \Delta = \tilde{\Delta})}{\partial^2 \Delta} \cdot \frac{\epsilon^2}{2}. \quad (3)$$

Equation (3) holds also with absolute values:

$$\Gamma(X, \epsilon) = \left| \text{MSE}(X, \Delta') - \text{MSE}(X, \tilde{\Delta}) \right| = \left| \frac{\partial^2 \text{MSE}(X, \Delta = \tilde{\Delta})}{\partial^2 \Delta} \cdot \frac{\epsilon^2}{2} \right|. \quad (4)$$

■

1.2

Lemma 2 Let X_U be a continuous random variable that is uniformly distributed in the interval $[-a, a]$. Assume that $Q_\Delta(X_U)$ is a uniform M -bit quantizer with a quantization step Δ . Then, the expected MSE is given as follows:

$$\text{MSE}(X_U, \Delta) = \frac{(a - 2^{M-1}\Delta)^3}{3a} + \frac{2^M \cdot \Delta^3}{24a}.$$

16 **Proof:** Given a finite quantization step size Δ and a finite range of quantization levels 2^M , the
 17 quantizer truncates input values larger than $2^{M-1}\Delta$ and smaller than $-2^{M-1}\Delta$. Hence, denoting by
 18 τ this threshold (i.e., $\tau \triangleq 2^{M-1}\Delta$), the quantizer can be modeled as follows:

$$Q_\Delta(x) = \begin{cases} \tau & x > \tau \\ \Delta \cdot \left\lfloor \frac{x}{\Delta} \right\rfloor & |x| \leq \tau \\ -\tau & x < -\tau \end{cases} \quad (5)$$

19 Therefore, by the law of total expectation, we know that

$$\begin{aligned} \mathbb{E} \left[(x - Q_\Delta(x))^2 \right] &= \\ \mathbb{E} \left[(x - \tau)^2 \mid x > \tau \right] \cdot P[x > \tau] &+ \\ \mathbb{E} \left[\left(x - \Delta \cdot \left\lfloor \frac{x}{\Delta} \right\rfloor \right)^2 \mid |x| \leq \tau \right] \cdot P[|x| \leq \tau] &+ \\ \mathbb{E} \left[(x + \tau)^2 \mid x < -\tau \right] \cdot P[x < -\tau] &. \end{aligned} \quad (6)$$

20 We now turn to evaluate the contribution of each term in Equation (6). We begin with the case of
 21 $x > \tau$, for which the probability density is uniform in the range $[\tau, a]$ and zero for $x > a$. Hence, the
 22 conditional expectation is given as follows:

$$\mathbb{E} \left[(x - \tau)^2 \mid x > \tau \right] = \int_{\tau}^a \frac{(x - \tau)^2}{a - \tau} \cdot dx = \frac{1}{3} \cdot (a - \tau)^2. \quad (7)$$

23 In addition, since x is uniformly distributed in the range $[-a, a]$, a random sampling from the interval
 24 $[\tau, a]$ happens with a probability

$$P[x > \tau] = \frac{a - \tau}{2a}. \quad (8)$$

25 Therefore, the first term in Equation (6) is stated as follows:

$$\mathbb{E} \left[(x - \tau)^2 \mid x > \tau \right] \cdot P[x > \tau] = \frac{(a - \tau)^3}{6a}. \quad (9)$$

26 Since x is symmetrical around zero, the first and last terms in Equation (6) are equal and their sum
 27 can be evaluated by multiplying Equation (9) by two.

28 We are left with the middle part of Equation (6) that considers the case of $|x| < \tau$. Note that the
 29 quantizer rounds input values to the nearest discrete value that is a multiple of the quantization step
 30 Δ . Hence, the quantization error, $e = x - \Delta \cdot \left\lfloor \frac{x}{\Delta} \right\rfloor$, is uniformly distributed and bounded in the range
 31 $[-\frac{\Delta}{2}, \frac{\Delta}{2}]$. Hence, we get that

$$\mathbb{E} \left[\left(x - \Delta \cdot \left\lfloor \frac{x}{\Delta} \right\rfloor \right)^2 \mid |x| \leq \tau \right] = \int_{-\frac{\Delta}{2}}^{\frac{\Delta}{2}} \frac{1}{\Delta} \cdot e^2 de = \frac{\Delta^2}{12}. \quad (10)$$

32 Finally, we are left to estimate $P[|x| \leq \tau]$, which is exactly the probability of sampling a uniform
 33 random variable from a range of 2τ out of a total range of $2a$:

$$P[|x| \leq \tau] = \frac{2\tau}{2a} = \frac{\tau}{a}. \quad (11)$$

34 By summing all terms of Equation (6) and substituting $\tau = 2^{M-1}\Delta$, we achieve the following
 35 expression for the expected MSE:

$$\mathbb{E} \left[(x - Q_\Delta(x))^2 \right] = \frac{(a - \tau)^3}{3a} + \frac{\tau}{a} \frac{\Delta^2}{12} = \frac{(a - 2^{M-1}\Delta)^3}{3a} + \frac{2^M \Delta^3}{24a}. \quad (12)$$

36 ■

1.3

Lemma 3 Let X_U be a continuous random variable that is uniformly distributed in the interval $[-a, a]$. Given an M -bit quantizer $Q_\Delta(X)$, the expected MSE $\mathbb{E}[(X - Q_\Delta(X))^2]$ is minimized by selecting the following quantization step size:

$$\tilde{\Delta} = \frac{2a}{2^M \pm 1} \approx \frac{2a}{2^M}. \quad (13)$$

Proof: We calculate the roots of the first order derivative of Equation (12) with respect to Δ as follows:

$$\frac{\partial \text{MSE}(X_U, \Delta)}{\partial \Delta} = \frac{1}{a} \left(2^{M-3} \Delta^2 - 2^{M-1} (a - 2^{M-1} \Delta)^2 \right) = 0. \quad (14)$$

Solving Equation (14) yields the following solution:

$$\tilde{\Delta} = \frac{2a}{2^M \pm 1} \approx \frac{2a}{2^M}. \quad (15)$$

■

2 Hyper parameters to reproduce the results in Section 5- Experiments

In the following section we describe the hyper parameters used in the experiments section. A fully reproducible code accompanies the paper.

2.1 Hyper parameters for Section 5.1- Robustness towards variations in quantization step size

In Table 1 we describe the hyper-parameters used in Fig. 4a and Fig. 4b in section 5.1 in the paper. We apply KURE on a pre-trained model from torch-vision repository and fine-tune it with the following hyper-parameters. When training phase ends we quantize the model using PTQ (Post Training Quantization) quantization method. All the other hyper-parameters like momentum and w-decay stay the same as in the pre-trained model.

Table 1: Hyper parameters for the experiments in section 5.1 - Robustness towards variations in quantization step size using PTQ methods

arch	kurtosis target (\mathcal{K}_T)	KURE coefficient (λ)	initial lr	lr schedule	batch size	epochs	fp32 accuracy
ResNet-50	1.8	1.0	1e-3	decays by a factor of 10 every 30 epochs	128	50	76.4

In Table 2 we describe the hyper-parameters used in Fig. 4c and Fig. 4d in section 5.1 in the paper. We combine KURE with QAT method during the training phase with the following hyper-parameters.

2.2 Hyper parameters for Section 5.2- Robustness towards variations in quantization bit-width

In Table 3 we describe the hyper-parameters used in Table 1 in section 5.2.1 in the paper. We apply KURE on a pre-trained model from torch-vision repository and fine-tune it with the following hyper-parameters.

In Table 4 we describe the hyper-parameters used in Fig. 5 in section 5.2.1 in the paper. We combine KURE with QAT method during the training phase with the following hyper-parameters.

Table 2: Hyper parameters for experiments in section 5.1 - Robustness towards variations in quantization step size using QAT methods

arch	QAT method	quantization settings (W/A)	kurtosis target (\mathcal{K}_T)	KURE coefficient (λ)	initial lr	lr schedule	batch size	epochs	acc
ResNet-18	DoReFa	4 / 4	1.8	1.0	1e-4	decays by a factor of 10 every 30 epochs	256	80	68.3
MobileNet-V2	DoReFa	4 / 8	1.8	1.0	5e-5	lr decay rate of 0.98 per epoch	128	10	66.9

Table 3: Hyper parameters for experiments in section 5.2 - Robustness towards variations in quantization bit-width using PTQ methods

architecture	kurtosis target (\mathcal{K}_T)	KURE coefficient (λ)	initial lr	lr schedule	batch size	epochs	fp32 accuracy
ResNet-18	1.8	1.0	0.001	decays by a factor of 10 every 30 epochs	256	83	70.3
ResNet-50					128	49	76.4
MobileNet-V2					256	83	71.3

Table 4: Hyper parameters for experiments in section 5.2 - Robustness towards variations in quantization bit-width using QAT methods

arch	QAT method	quantization settings (W/A)	kurtosis target (\mathcal{K}_T)	KURE coefficient (λ)	initial lr	lr schedule	batch size	epochs	acc
ResNet-18	LSQ	6 / 6	1.8	1.0	1e-3	decays by a factor of 10 every 20 epochs	128	60	70.1
ResNet-18	LSQ	4 / 4					128	60	69.3
ResNet-50	LSQ	6 / 8					64	50	76.5

3 Robustness towards variations in quantization bit-width- additional results

In Fig. 5 in the paper we demonstrated robustness to variations in quantization bit-width of QAT models. we used LSQ method as our QAT model. In Fig. 1 we demonstrate the improved robustness with different QAT methods (DoReFa and LSQ) and ImageNet models.

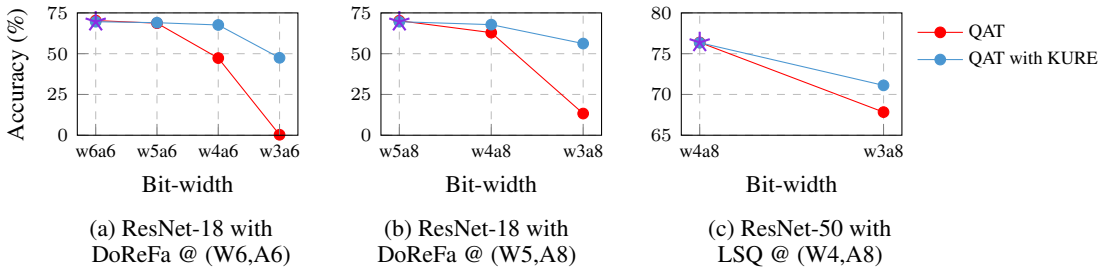


Figure 1: Bit-width robustness comparison of QAT model with and without KURE on different ImageNet architectures. The \star is the original point to which the QAT model was trained.

69 **4 Statistical significance of results on ResNet-18/ImageNet trained with**
70 **DoReFa and KURE**

Table 5: Mean and standard deviation over multiple runs of ResNet-18 trained with DoReFa and KURE

architecture	QAT method	quantization settings (W/A)	Runs	Accuracy, % (mean \pm std)
ResNet-18	DoReFa	4 / 4	3	(68.4 \pm 0.09)




Figure 2: The network has been trained for quantization step size $\tilde{\Delta}$. Still, the quantizer uses a slightly different step size Δ . Small changes in optimal step size $\tilde{\Delta}$ cause severe accuracy degradation in the quantized model. KURE significantly enhances the model robustness by promoting solutions that are more robust to uncertainties in the quantizer design (ResNet-18 on ImageNet).